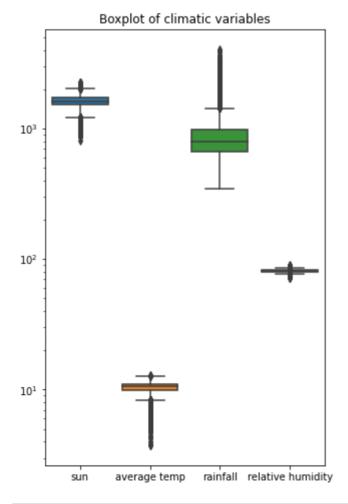
```
In []:
         import pandas as pd
         from sklearn.linear model import LinearRegression
         from sklearn.model selection import train test split, GridSearchCV
         from sklearn import metrics
         from sklearn import tree
         from sklearn.ensemble import ExtraTreesRegressor
         from sklearn.neighbors import KNeighborsClassifier
         from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScale
         import matplotlib.pyplot as plt
         import seaborn as sns
         import numpy as np
         from imblearn import under sampling
In []:
         #loading intermediate dataset ukbms and haduk, created in HADUK intermedia
         df = pd.read csv('BioD year site latlong.csv')
         df.drop(columns=df.columns[0], inplace = True)
In []:
         df1 climate = df.drop(columns=['Simpsons Index','SPECIES RICHNESS','TOTAL S
         df1 climate.describe()
         fig,ax = plt.subplots(figsize=(5,8))
         bxplot = sns.boxplot(data=df1 climate)
         bxplot.set yscale('log')
         plt.title('Boxplot of climatic variables')
```

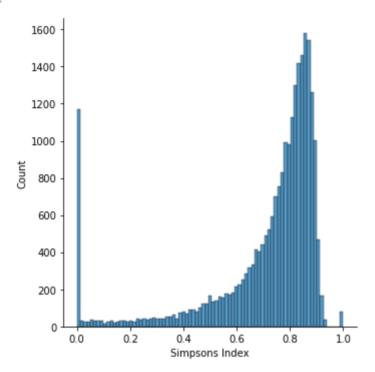
Out[]: Text(0.5, 1.0, 'Boxplot of climatic variables')

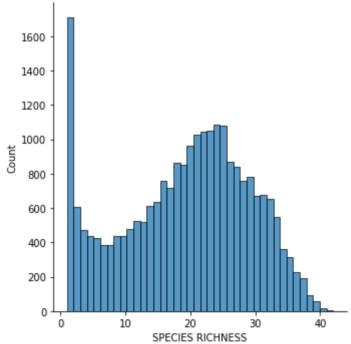


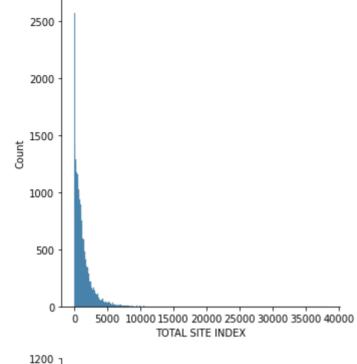
```
In [ ]: sns.displot(df['Simpsons Index'])
```

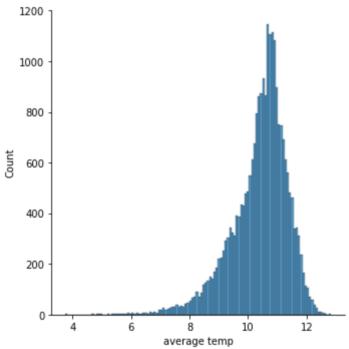
```
sns.displot(df['SPECIES RICHNESS'])
sns.displot(df['TOTAL SITE INDEX'])
sns.displot(df['average temp'])
sns.displot(df['sun'])
sns.displot(df['rainfall'])
sns.displot(df['relative humidity'])
```

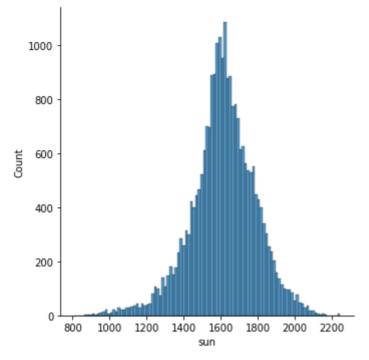
Out[]: <seaborn.axisgrid.FacetGrid at 0x7fc2a6c1b220>

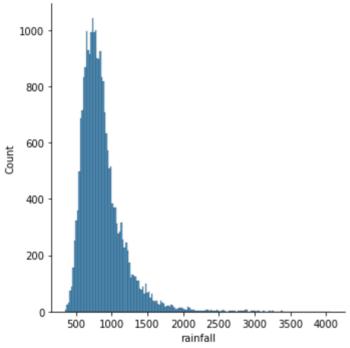


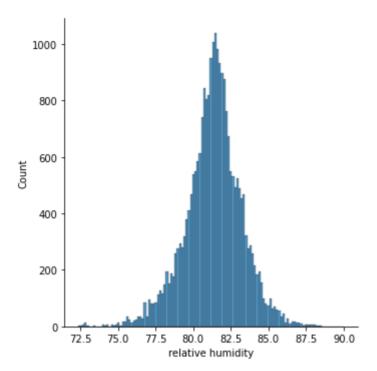






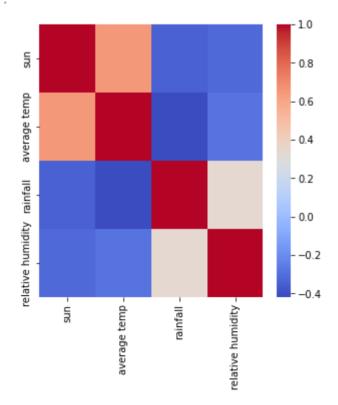






Multiple Linear Regression, carried out on climatic variables (x) and species richness (y):

Out[]: <AxesSubplot:>



```
In []: x_train.drop(columns=['sun'],inplace=True)
    x_test.drop(columns=['sun'],inplace=True)
```

```
In [ ]:
         #training multiple linear regression model
         model = LinearRegression()
         model.fit(x_train, np.log(y_train))
         print(model.coef )
         print(model.intercept_)
         [[ 0.06703855 -0.00027393 -0.04370613]]
         [5.80492544]
In []:
         preds SR = model.predict(x test)
In [ ]:
         #plotting residuals
         residuals = np.log(y_test)-preds_SR
         plt.scatter(preds_SR,residuals,5)
         plt.ylabel('Residuals')
         plt.xlabel('Predicted values')
         plt.title('Residuals plot')
         plt.show()
                              Residuals plot
            1
        Residuals
           -2
           -3
                1.8
                      2.0
                           22
                                24
                                     2.6
                                           2.8
                                                3.0
                                                     3.2
                              Predicted values
In [ ]:
         mse = metrics.mean_squared_error(np.log(y_test),preds_SR)
         r2 = metrics.r2_score(np.log(y_test),preds_SR)
         print(mse)
         print(np.sqrt(mse))
         print(r2)
        0.7450615386919197
        0.8631694727525526
        0.04805919630695088
In []:
         #Checking diversity
         x = df[['average temp', 'rainfall', 'relative humidity']]
         y= df[['Simpsons Index']]
         x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3, random_
         model1 = LinearRegression()
         model1.fit(x_train, y_train)
         print(model1.coef_)
         print(model1.intercept )
         preds_Div = model1.predict(x_test)
         mse = metrics.mean_squared_error(y_test,preds_Div)
         r2 = metrics.r2_score(y_test,preds_Div)
         print('mse: ' + str(mse))
```

```
print('rmse: ' + str(np.sqrt(mse)))
print('r2: ' + str(r2))
```

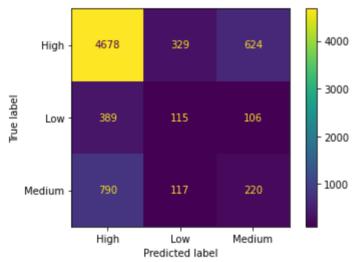
```
[[ 9.50157468e-03 -7.94726070e-05 -8.92072927e-03]]
[1.40639564]
mse: 0.04964239506285567
rmse: 0.22280573390928626
r2: 0.03478576404689471
```

Decision tree classifier:

```
In []: #multiclass classifcation decision tree classifier
    x_c = df[['sun', 'average temp', 'rainfall', 'relative humidity']]
    y_c = df[['Diversity']]
    x_ctrain,x_ctest,y_ctrain,y_ctest = train_test_split(x_c,y_c,test_size=0.3,
    dtree_model = tree.DecisionTreeClassifier()
    dtree_model.fit(x_ctrain,y_ctrain)
    dtree_preds = dtree_model.predict(x_ctest)

cm = metrics.confusion_matrix(y_ctest,dtree_preds)
    acc = metrics.accuracy_score(y_ctest, dtree_preds)
    metrics.ConfusionMatrixDisplay.from_predictions(y_ctest,dtree_preds)
    print(acc)
```

0.6803745928338762



```
In []: #d tree classifier tunin
    ctree_tuning = tree.DecisionTreeClassifier()
    c_tuning_model=GridSearchCV(ctree_tuning,param_grid=parameters,scoring='acc
    tuning_sample = df.sample(frac=0.3) #sampling dataset to prevent crashses w
    x_ctuning = tuning_sample[['sun','average temp','relative humidity','rainfa
    y_ctuning = tuning_sample[['Diversity']]
```

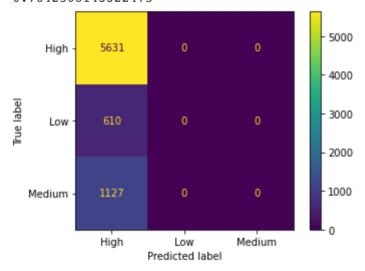
```
In [ ]: c_tuning_model.fit(x_ctuning,y_ctuning)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package s/sklearn/tree/_classes.py:298: FutureWarning: `max_features='auto'` has bee n deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features='sqrt'`.

metrics.ConfusionMatrixDisplay.from predictions(y ctest,tuned ctree preds)

warnings.warn(0.7642508143322475

print(acc)



KNN CLASSIFICATION:

```
In []: #multiclass classifcation k-nearest neighbours
knn_model = KNeighborsClassifier(n_neighbors = 10).fit(x_ctrain,y_ctrain)
#accuracy = knn_model.score(x_ctest,y_ctest)

knn_preds = knn_model.predict(x_ctest)
cm_knn = metrics.confusion_matrix(y_ctest,knn_preds)
acc_knn =metrics.accuracy_score(y_ctest,knn_preds)

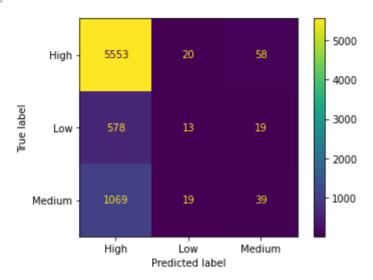
metrics.ConfusionMatrixDisplay.from_predictions(y_ctest,knn_preds)
print(acc_knn)
metrics.fl_score(y_ctest,knn_preds, average='weighted')
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package s/sklearn/neighbors/_classification.py:207: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self._fit(X, y)
0.7607220412595005

Out[]:

0.6743548782062548



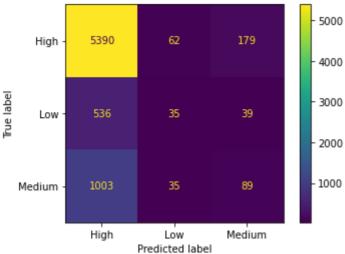
KNN SCALED X VARIABLES:

```
In [ ]:
         #scaled knn
         scaler = StandardScaler()
         df scale = df.copy()
         features = [['sun','average temp','rainfall','relative humidity']]
         for feature in features:
             df scale[feature] = scaler.fit transform(df scale[feature])
         x_scaled = df_scale[['sun','average temp','rainfall','relative humidity']]
         y s = df scale[['Diversity']]
         x_strain,x_stest,y_strain,y_stest = train_test_split(x_scaled,y_s,test_size
         knn scaled model = KNeighborsClassifier()
         knn scaled model.fit(x strain,y strain)
         knn_scaled_preds = knn_scaled_model.predict(x_stest)
         acc scaled = metrics.accuracy score(y stest,knn scaled preds)
         metrics.ConfusionMatrixDisplay.from predictions(y stest,knn scaled preds)
         print(acc scaled)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package s/sklearn/neighbors/_classification.py:207: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self. fit(X, y)

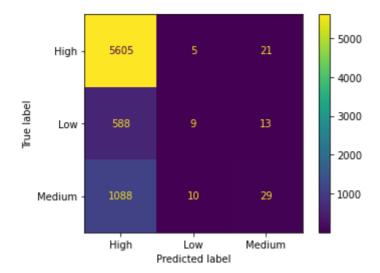
```
return self._fit(x, y)
0.748371335504886
```



return self._fit(X, y)

0.7658794788273615
0.6733164631286662

```
In [ ]:
         metrics.fl score(y stest,knn scaled preds, average='weighted')
        0.682738313321154
Out[ ]:
In [ ]:
         #knn hyperparameter tuning
         n neighbors = [5,10,15,20,25,30]
         leaf size= list(range(1,25))
         p=[1,2]
         knn_hyperparameters = dict(n_neighbors=n_neighbors,leaf_size=leaf_size,p=p)
         knn tuned = KNeighborsClassifier()
         knn_tuned_model = GridSearchCV(knn_tuned,knn_hyperparameters, cv=3)
In []:
         knn_tuned_model.fit(x_strain,y_strain.values.ravel())
         print('best leaf size:', knn tuned model.best estimator .get params()['leaf
         print('best p:', knn_tuned_model.best_estimator_.get_params()['p'])
         print('best n neighbors:', knn tuned model.best estimator .get params()['n
        best leaf size: 2
        best p: 1
        best n_neighbors: 30
In []:
         tuned knn model = KNeighborsClassifier(leaf size=2,p=1,n neighbors=30)
         tuned knn model.fit(x strain,y strain)
         tuned_knn_preds = tuned_knn_model.predict(x_stest)
         acc tuned knn = metrics.accuracy score(y stest,tuned knn preds)
         metrics.ConfusionMatrixDisplay.from predictions(y stest, tuned knn preds)
         print(acc tuned knn)
         print(metrics.fl_score(y_stest,tuned_knn_preds, average='weighted'))
        /Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package
        s/sklearn/neighbors/_classification.py:207: DataConversionWarning: A column-
        vector y was passed when a 1d array was expected. Please change the shape of
        y to (n_samples,), for example using ravel().
```



Condensed Neareset Neighbor Undersamling:

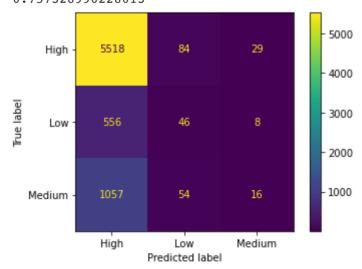
```
In []:
    undersample = under_sampling.CondensedNearestNeighbour()
    x_under,y_under = undersample.fit_resample(x_strain,y_strain)

In []:
    cnn_knn = KNeighborsClassifier(leaf_size=2,p=1,n_neighbors=30)
    cnn_knn.fit(x_under,y_under)
    cnn_knn_preds = cnn_knn.predict(x_stest)

    acc_knn_cnn = metrics.accuracy_score(y_stest,cnn_knn_preds)
    metrics.ConfusionMatrixDisplay.from_predictions(y_stest,cnn_knn_preds)
    print(acc_knn_cnn)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package s/sklearn/neighbors/_classification.py:207: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self._fit(X, y)
0.757328990228013



```
In [ ]:
    rus_undersample = under_sampling.RandomUnderSampler(random_state=44,replace
    x_rus,y_rus = undersample.fit_resample(x_strain,y_strain)
```

```
In [ ]:
    rus_knn = KNeighborsClassifier(leaf_size=2,p=1,n_neighbors=30)
    rus_knn.fit(x_rus,y_rus)
```

```
rus_knn_preds = rus_knn.predict(x_stest)

acc_knn_rus = metrics.accuracy_score(y_stest,rus_knn_preds)
metrics.ConfusionMatrixDisplay.from_predictions(y_stest,rus_knn_preds)
print(acc_knn_rus)
```

/Library/Frameworks/Python.framework/Versions/3.9/lib/python3.9/site-package s/sklearn/neighbors/_classification.py:207: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

return self._fit(X, y)

0.756786102062975

